

Cover Letter -- IEEE TCSVT Submission

To: Editor-in-Chief, IEEE Transactions on Circuits and Systems for Video Technology (TCSVT)

From: Qiming Bao (joint first author and first corresponding author) and Yanbing Bai (second corresponding author), on behalf of all co-authors

Qiming Bao -- The University of Auckland, Auckland 1010, New Zealand --
qiming.bao@auckland.ac.nz

Yanbing Bai -- Center for Applied Statistics, School of Statistics, Renmin University of China, Beijing 100872, China -- ybbai@ruc.edu.cn

Authorship: Jing An and Qiming Bao are joint first authors (equal contribution); Qiming Bao and Yanbing Bai are joint corresponding authors.

Date: [INSERT SUBMISSION DATE]

Manuscript title: *An End-to-End Multimodal System for Subtitle Recognition and Chinese-Japanese Translation with Speech Synthesis in Short Dramas*

Dear Editor,

We are pleased to submit our manuscript, listed above, for consideration as a Regular Paper in *IEEE Transactions on Circuits and Systems for Video Technology*. The paper presents what is, to the best of our knowledge, the first end-to-end multimodal localization system specifically designed for Chinese short-form drama (*duanju*), spanning the three tightly coupled video-processing stages of visual subtitle recognition (frame-level OCR), audio transcription with adaptive multimodal fusion, neural machine translation, and zero-shot Japanese voice cloning for dubbed audio.

Why TCSVT is the appropriate venue

Short-form drama localization is fundamentally a circuits-and-systems-for-video problem: the subtitle text, the audio waveform, and the visual frame are jointly constrained, and progress on any single modality in isolation is insufficient for a deployable video pipeline. Our paper contributes (i) a confidence-adaptive OCR/ASR fusion algorithm that gates between modalities using ASR log-probabilities; (ii) a domain-adapted vision-language translation model that exploits visual context for register-aware Chinese-Japanese rendering; (iii) a controlled comparison of three zero-shot Japanese voice cloning systems with a reference-duration ablation; and (iv) a manually annotated 79-episode benchmark with end-to-end real-time-factor profiling. These contributions sit squarely within the TCSVT scope of video signal analysis, video-understanding systems, multimodal video processing, and applied learning for video content, and the work is unlikely to find a natural home at a single-modality venue (vision, NLP, or speech).

Relation to prior conference publication

This manuscript is a substantial extension of our preliminary conference paper accepted to **IEEE ICASSP 2026** (Session MMSP-P23), which introduced an early version of the OCR+ASR fusion idea and a small-scale LoRA fine-tuned Qwen2.5-3B translation model on a subset of the data. The present journal version extends that work along six independent dimensions, contributing well beyond the new-material expectation under the IEEE Policy on Prior Publication for journal extensions of conference papers. A separate **Summary of Differences** document detailing the extensions has been uploaded with this submission as required by IEEE policy.

1. **Comprehensive recognition benchmark** (Section 5.3): seven OCR backbones evaluated under matched conditions (Qwen3-VL, Qwen2-VL, InternVL2, GOT-OCR2.0, RapidOCR, EasyOCR, TrOCR) versus the single Qwen2-VL configuration in the conference version.
2. **Frame-rate ablation** (Section 5.4): a new study quantifying the 1 fps -> 5 fps recall gap from 27.7% -> 100%, absent from the conference paper.
3. **Upgraded fusion algorithm** (Section 4.2.3, Algorithm 1): replaces the conference paper's fixed similarity threshold with a confidence-adaptive gate using Whisper's segment log-probabilities, formalized as new pseudocode.
4. **Substantially stronger translation results**: LoRA fine-tuned **Qwen3-VL-4B** achieves **BLEU 29.99 / chrF++ 39.34 / COMET 0.870**, versus the BLEU 9.84 reported in the conference paper. We additionally benchmark NLLB-200 and Qwen2.5-7B baselines.
5. **Complete speech-synthesis stage** (Sections 4.4 and 5.7): GPT-SoVITS, F5-TTS, and EdgeTTS comparison with reference-duration ablation, **entirely new** relative to the conference paper.
6. **New analyses**: visual-context translation ablation (+3.35 BLEU), computational cost / real-time-factor profiling, and a 100-segment qualitative error taxonomy.

Originality and ethical compliance

- This manuscript is original work; it has not been previously published, and it is not under consideration at any other journal or conference.
- All listed co-authors have read and approved the submission and have agreed on author order.
- The dataset was provided by a commercial entertainment company under a research-use agreement (Section *Ethics Statement*). All performers consented to broadcast distribution; human annotators participated voluntarily and were compensated.
- We acknowledge the dual-use risk of zero-shot voice cloning and have chosen not to release reference-audio-to-cloned-voice mappings; only short, content-neutral synthetic samples will be released for inspection (see Supplementary Materials).
- All authors declare **no competing financial or non-financial interests** related to this work.

Suggested reviewers (optional)

We do not propose specific reviewers, but we suggest the editorial board consider experts in any of the following overlapping areas: (i) video subtitling and multimedia localization, (ii) multimodal fusion for speech and video understanding, (iii) low-resource and domain-adapted neural machine translation for video content, and (iv) zero-shot voice cloning and TTS evaluation for video dubbing.

Reviewer-excluded list

We respectfully request that the following individuals not be invited as reviewers due to recent collaboration: *[INSERT NAMES IF ANY; LEAVE EMPTY OTHERWISE]*.

Compliance with TCSVT submission requirements

- **Format:** IEEE IEEEtran class, journal option (two-column, single-blind).
- **References:** IEEE Reference Format (IEEEtran.bst).
- **Keywords:** IEEE keywords block on the first page.
- **Summary of differences:** uploaded as a separate document for the journal extension of our ICASSP 2026 paper, per IEEE Policy on Prior Publication.
- **Supplementary materials:** README, configuration files, per-episode metric breakdowns, translation example pairs, and short anonymized synthesis samples are provided in the supplementary archive.
- **Reproducibility, Data Availability, and Ethics statements:** provided as separate sections at the end of the manuscript.

We thank you and the reviewers in advance for considering our submission and look forward to your feedback.

Sincerely,

Qiming Bao -- first corresponding author

The University of Auckland, Auckland 1010, New Zealand

qiming.bao@auckland.ac.nz

Yanbing Bai -- second corresponding author

Center for Applied Statistics, School of Statistics, Renmin University of China

59 Zhongguancun Street, Haidian District, Beijing 100872, China

ybbai@ruc.edu.cn

on behalf of all co-authors: Jing An (joint first author), Rui-Yang Ju, Hao-fei Chang, Jinhua Su, and Xin Qu